

Building a Semantic Search Engine (workshop)



SCAN ME

Outline

— — —

1. Sparse Embedding
2. Dense Embedding
3. Approximate Nearest Neighbor (ANN)
5. Production Considerations

About Us

— — —



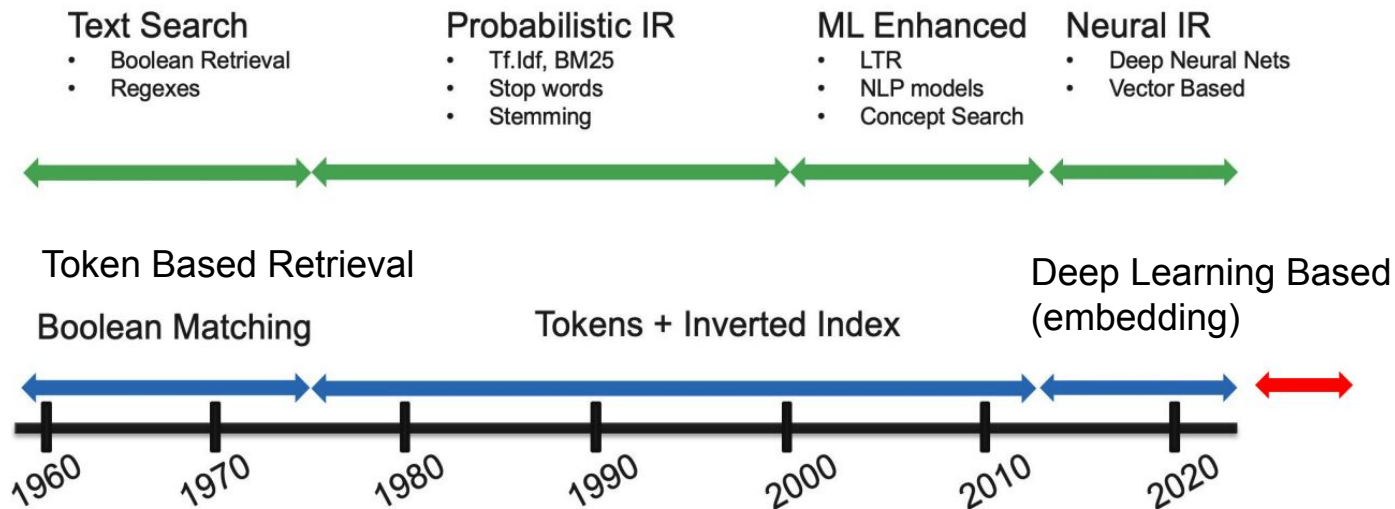
Ravi Yadav
[Linkedin](#)



Nidhin Pattaniyil
[Linkedin](#)

Walmart: ML Engineers on the Search team

History of Information Retrieval



Hughes, Simon. "Semantic Product Search – Vector Search for E-Commerce." Conference Presentation at Haystack 2021, https://haystackconf.com/files/slides/haystack2021/Hughes-Haystack_2021_Semantic_Product_Search.pdf, September 29, 2021.

Token based retrieval

Boolean Retrieval

- Queries and documents are represented as bag of words
- Query terms are connected with boolean operators
- Disadvantages:
 - Filtering more than retrieval
 - Terms have same weights
 - No ranking

pictures of kitten playing 

Query Processing (tokenize , stop words, stemming)

["pictures", "of", "kitten", "playing"]

["picture", "kitten", "play"]

picture OR kitten OR play

Which Document is most relevant ?

images of cat playing

video of kitten ..
having fun

kitten sleeping
dog playing in park
.....

IG cooking pictures

score: 1

score: 1

score: 3

TF-IDF

TF-IDF is a measure of originality of a word by comparing the number of times a word appears in a doc with the number of docs the word appears in.

$$\text{TF-IDF} = \text{TF}(t, d) \times \text{IDF}(t)$$


Term frequency

Number of times term t appears in a doc, d

Inverse document frequency


$$\log \frac{1 + \overset{\text{\# of documents}}{n}}{1 + \underset{\text{Document frequency of the term } t}{df(d, t)}} + 1$$

BM25



idf - how popular
is the term in the
corpus?

$$\text{bm25}(d) = \sum_{t \in q, f_{t,d} > 0} \log \left(1 + \frac{N - df_t + 0.5}{df_t + 0.5} \right) \cdot \frac{f_{t,d}}{f_{t,d} + k \cdot \left(1 - b + b \frac{l(d)}{\text{avgdl}} \right)}$$




saturation
curve - limit
influence of tf
on the score



length weighing -
tweak influence of
document length

73



ElasticSearch

— — —

- Open source search engine based on Lucene library
- Supports BM25 and other similarities ([link](#))
- Supports boosting , filtering , phrase match, autocomplete
- Distributed : index shards, replicas



elasticsearch

ElasticSearch Schema

```
{
  "mappings": {
    "properties": {
      "title": {
        "type": "text"
      },
      "description": {
        "type": "text"
      },
      "brand": {
        "type": "keyword"
      },
      "product_type": {
        "type": "keyword"
      },
      "price": {
        "type": "double"
      }
    }
  }
}
```

full text field

regular text field

numeric field

Possible schema for an e-commerce item

ElasticSearch Query

Query: Nike shoes under 100\$

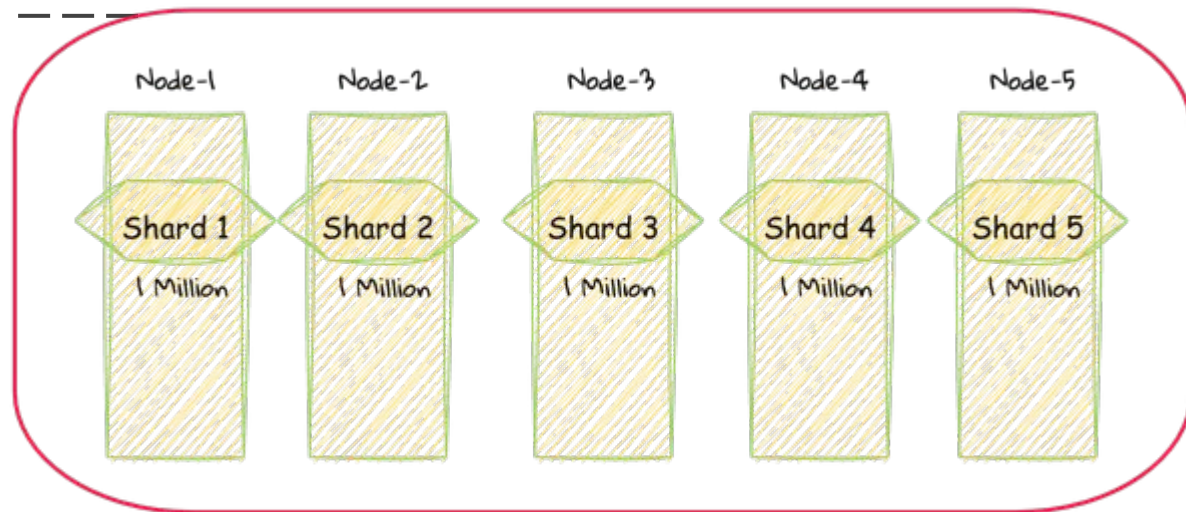
```
{
  "query": {
    "multi_match": {
      "query": "Nike shoe under 100$",
      "fields": ["title^2", "Description^1"]
    },
    "bool": {
      "filter": [
        { "term": { "brand": "nike" } }
      ]
    },
    "filtered": {
      "filter": {
        "range": {
          "price": {
            "lte": 100
          }
        }
      }
    }
  }
}
```

Search for query tokens in title and description
weight matches found in title double

filter items who has brand nike

filter items with price <= 100

Index with multiple shards

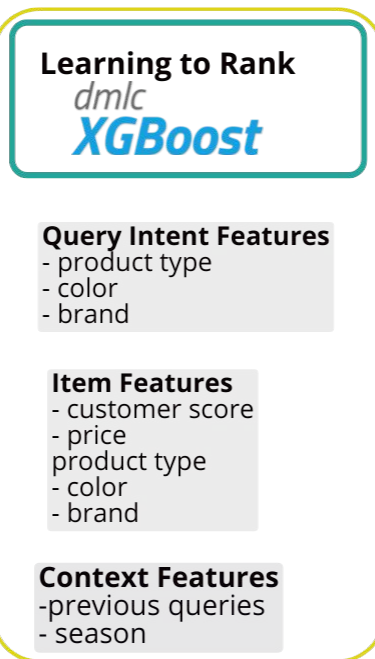



- Index can be composed of shards
- Reading / Writing can be faster

Learning to Rank

Red Dress

Original Order



New Order



Issues with Sparse Representation

— — —

- Lexical GAP: Covid vs Coronavirus vs Omicron variant
- Ambiguity: bank (institution) vs bank (geography)
- Position matters: “river bank” vs “bank river”
- Lack of Contextualized embeddings

She will *park* the car so we can walk in the *park*.



Lab

Jupyter Hub: <https://hub.np.training>

Repo: <https://bit.ly/search-workshop-2022>

Lab 1 Goals

- Explore tokenization and some preprocessing
- Building a simple in-memory retrieval system using BM-25 (elastic search)

Query: Two dogs playing in the snow

— — —



Photo title: two white dogs

Photo by [Konstantin Pudan](#) on [Unsplash](#)

Distance: 14.775831



Photo title: Two baby elephants playing

Photo by [Jason Zhao](#) on [Unsplash](#)

Distance: 14.002172

Query: boy and girl on a beach

— — —




Photo title: A drone shot of a girl laying on Pink Beach in Indonesia

Photo by [Marcus Woodbridge](#) on [Unsplash](#)

Distance: 12.359471




Photo title: Beautiful girl on a sunflower field

Photo by [Marina Montoya](#) on [Unsplash](#)

Distance: 11.130171

Query: light at the end of the tunnel

— — —



Photo title: At the end of the tunnel

Photo by [Dejan Zakic](#) on [Unsplash](#)

Distance: 25.557308



Photo title: Beach at the end of a day

Photo by [Angelo Pantazis](#) on [Unsplash](#)

Distance: 15.358174

Embedding Based Retrieval

Dense Embeddings

Word Representation

car: [0.2 , 0.3, 0.7]

automobile: [0.2 , 0.3, 0.7]

Similar concepts have similar embeddings


Regardless of content length, similar items should have similar embeddings

Size of embedding is independent of #tokens

Review Representation

Review 1: 🍦 was great [0.5 , 0.1, ... , 0.6]


Review 2: Chocolate ice cream was the best .. [0.5 , 0.1, ... , 0.5]



Word2Vec


My *clever* *cute* *kitten* is playing with

window size = 2



- Published in 2013
- Represent each word as a dense vector
- Uses a neural network model to capture linguistic concept of words

Instagram: ig2Vec

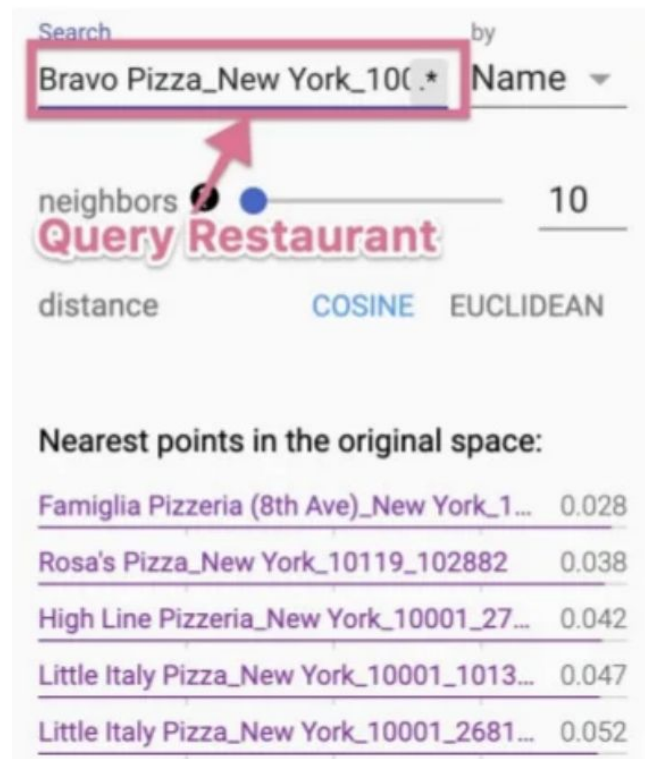



Medvedev, Ivan, Haotian Wu, and Taylor Gordon. "Powered by AI: Instagram's Explore Recommender System." Powered by AI: Instagram's Explore recommender system, November 2019.

<https://ai.facebook.com/blog/powered-by-ai-instagrams-explore-recommender-system/>

Grubhub: Rest2Vec

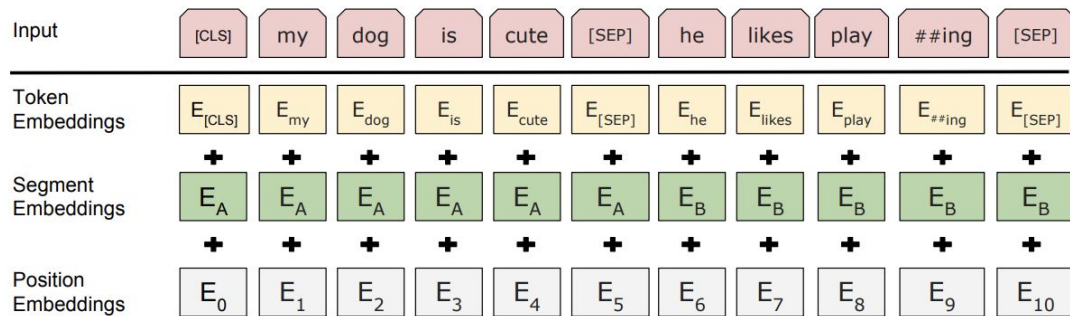
— — —



Content from Pydata Talk "Alex Egg, Emily A Ray, Parin Choganwala: Discover your latent food graph with this 1 weird trick | PyData New York 2019"

https://www.youtube.com/watch?v=aRUaEt1q7BY&t=3s&ab_channel=PyData

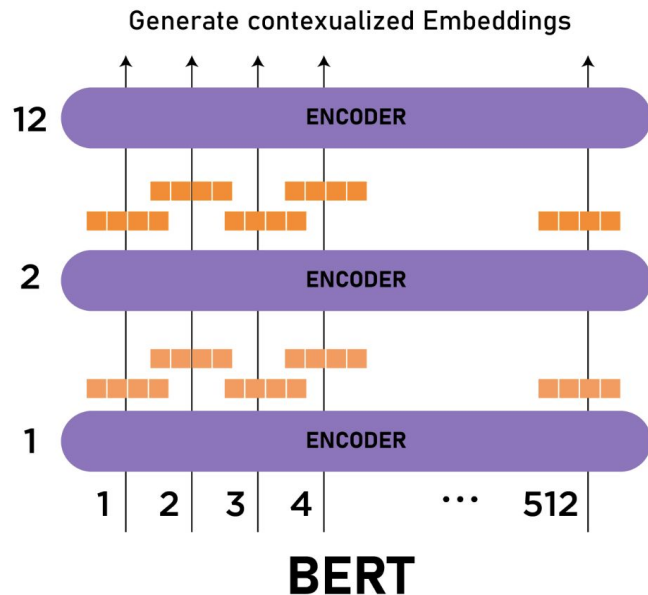
Bidirectional Encoder Representations from Transformers (BERT)




Input to the model, contains token , segment and position embedding

Image from BERT paper <https://arxiv.org/pdf/1810.04805.pdf>

BERT




The output of each encoder layer can be used to represent the feature for that token



- BERT is composed of multiple encoders
- There is embedding for each token after each encoder
- Sentence embedding can be created from pooling
- Pooled embeddings are not best for similarity search

Sentence Pair Encoder - Cross-encoder




Pros

- Accurate classification

Cons

- Need new encoding for each pair
- Computationally inefficient for information retrieval.

Sentence Pair Encoder - Bi-Encoder



Pros

- Each sentence is encoded separately
- Distance between two sentence embeddings can be measured
- Faster and more computationally efficient

Cons

- Bi-encoders performance is poor compared to cross-encoders for supervised learning problems



Lab

Jupyter Hub: <https://hub.np.training>

Repo: <https://bit.ly/search-workshop-2022>

Lab 2 Goals

— — —

- Explore Text Embedding model
- ANN Retrieval using Milvus

Query: Two dogs playing in the snow

— — —



Photo title: brown and black dogs running on snow

Photo by [Luka Vovk](#) on [Unsplash](#)

Distance: 0.7353920340538025



Photo title: tan dog playing on snow

Photo by [Hitter Rudolf](#) on [Unsplash](#)

Distance: 0.6460988521575928

Query: boy and girl on a beach

— — —




Photo title: children enjoying the beach

Photo by [Daria Nepriakhina](#) on [Unsplash](#)

Distance: 0.7077873945236206



Photo title: Couple on the beach

Photo by [Scott Webb](#) on [Unsplash](#)

Distance: 0.6981672048568726

Query: light at the end of the tunnel

— — —



Photo title: Light Keyhole

Photo by [Leon LEE](#) on [Unsplash](#)

Distance: 0.5932173728942871





Photo title: light direction

Photo by [Zhang Kenny](#) on [Unsplash](#)

Distance: 0.5392670035362244

Different types of negatives

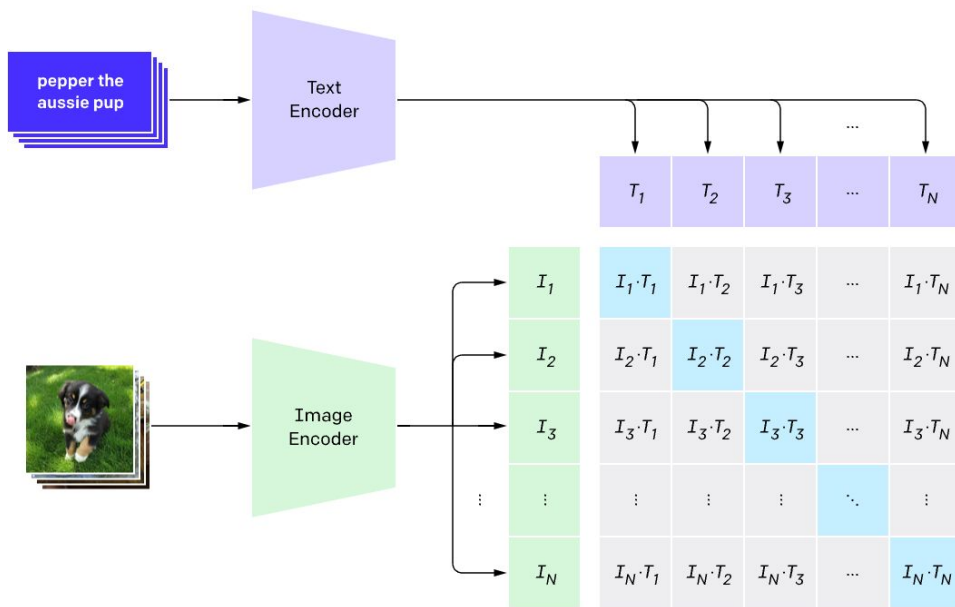
Anchor, Positive Ex, Negative Ex



Multiple Modality: Vision Transformers (ViT)

— — —

- Vision Transformer models like CLIP model uses two encoders (text and image).
- These two models are trained in parallel and optimized via a contrastive loss function
- End result is where you can search by text or image



Encoder can include multiple features

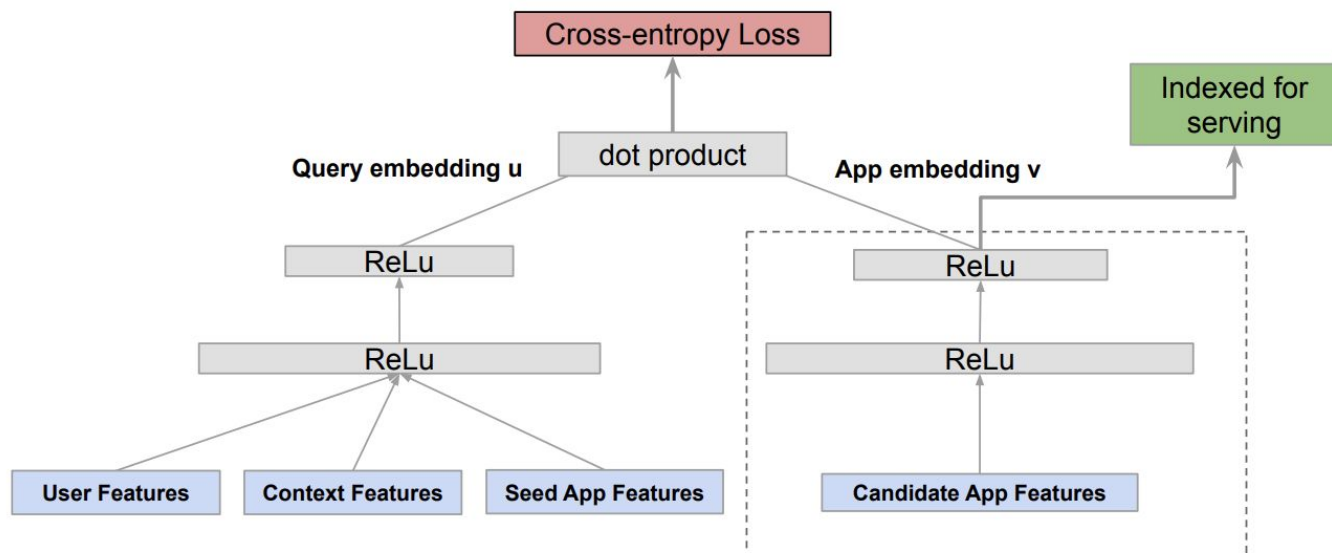


Figure 5: Two-tower model architecture for Google Play app recommendation.



Lab

Jupyter Hub: <https://hub.np.training>

Repo: <https://bit.ly/search-workshop-2022>

Lab 3 Goals

— — —

- Building a simple in-memory retrieval system using a multi-modal model like CLIP

Query: Two dogs playing in the snow

— — —



Photo title: brown and black dogs running on snow

Photo by [Luka Vovk](#) on [Unsplash](#)

Distance: 0.31937479502958604



Photo title: Friends will be friends

Photo by [Tadeusz Lakota](#) on [Unsplash](#)

Distance: 0.309985660979709

Query: boy and girl on a beach

— — —



Photo title: White sands

Photo by [Toa Heftiba](#) on [Unsplash](#)

Distance: 0.3079377628510268




Photo title: children enjoying the beach

Photo by [Daria Nepriakhina](#) on [Unsplash](#)

Distance: 0.30199767578628933

Query: light at the end of the tunnel

— — —



Photo title: Gateway to heaven

Photo by [Niklas Schweinzer](#) on [Unsplash](#)

Distance: 0.2948713358039129




Photo title: cave during golden hour

Photo by [kiwi thompson](#) on [Unsplash](#)


Distance: 0.292107594402263

Approximate Nearest Neighbors




In a flat index / no index, our query vector is compared against all the vectors in the database.

Inverted File Index: Building




- Find centroid and create Voronoi Cells
- Number of centroids is determined by **nlist** parameter

Inverted File Index: Building



- Built Voronoi Index with 15 centroids
- Memory usage is not reduced
- But retrieval is faster


Inverted File Index: Searching



nprobe = 1


- Find the distance between query vector and all the centroids
- Query all the elements inside the n closest cluster determined by **nprobe** parameter

Inverted File Index: Searching



- Increasing the value of **nprobe**, improves recall but increases latency
- If $nprobe = nlist$, similar to flat index


Hierarchical Navigable Small Worlds (HNSW)



The search process through the multi-layer structure of an HNSW graph.

- Hierarchical Graph
- Every layer stores neighbors
- Earlier layers are sparse
- Common implementation in most production ANN databases

ANN Benchmarks



Source:
[ANN Benchmarks](#)
[Subset](#)

Lot of ANN Option

— — —

FAISS




Evaluation Considerations

— — —

- Managed vs Self-hosted
- Performance
- Update Embeddings / Partial Updates
- Metadata Filtering
- Filtering / Hybrid Retrieval
- Plugins

Filtering



Closest 10
candidates don't
meet our filter

Hybrid / Full Retrieval

Supports



elastic




vespa

- In ES, **disjunction** of knn and bm25 match.
- The score of each hit is the sum of the knn and query scores. A boost can be specified

```
POST image-index/_search
{
  "query": {
    "match": {
      "title": {
        "query": "mountain lake",
        "boost": 0.9
      }
    }
  },
  "knn": {
    "field": "image-vector",
    "query_vector": [54, 10, -2],
    "k": 5,
    "num_candidates": 50,
    "boost": 0.1
  },
  "size": 10
}
```

ElasticSearch example of Hybrid Retrieval

Ex from Elastic Doc [link](#)



Lab

Jupyter Hub: <https://hub.np.training>

Repo: <https://bit.ly/search-workshop-2022>

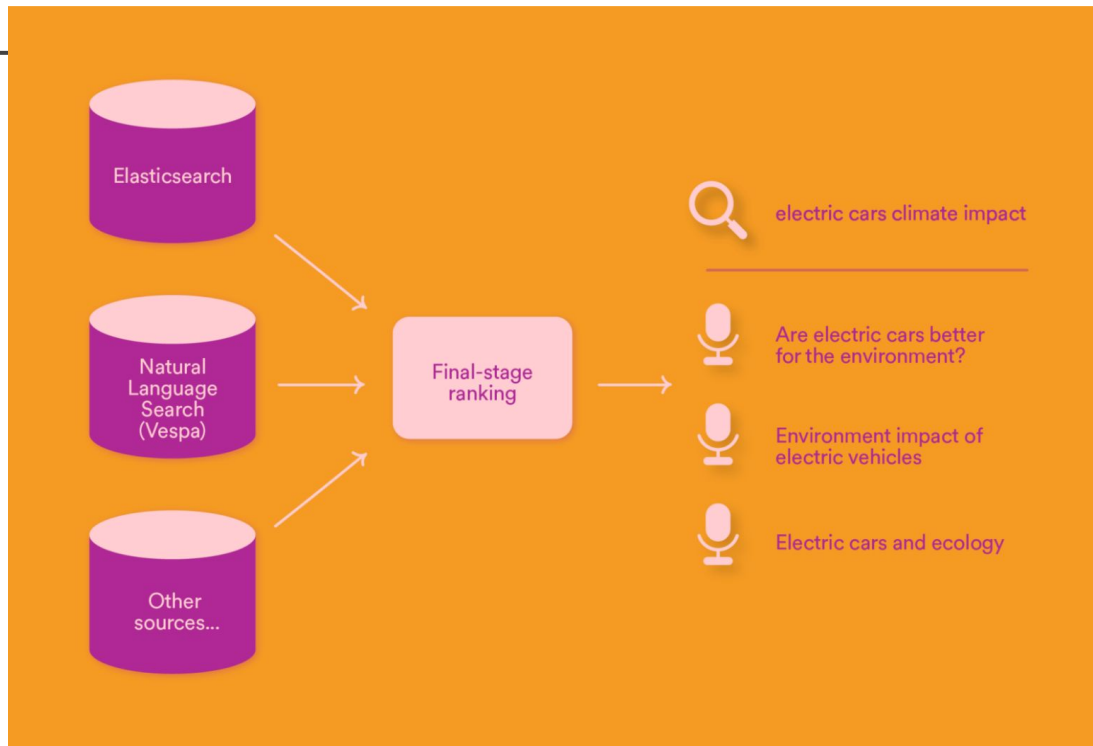
Lab 4 Goals

— — —

- Build a flat ANN index
- Optimize retrieval by building an IVF Index

Conclusion

Multi Source Retrieval



Spotify Podcast Episode Search

Retrieval from ES + ANN + other sources

Multi Stage Retrieval and Ranking

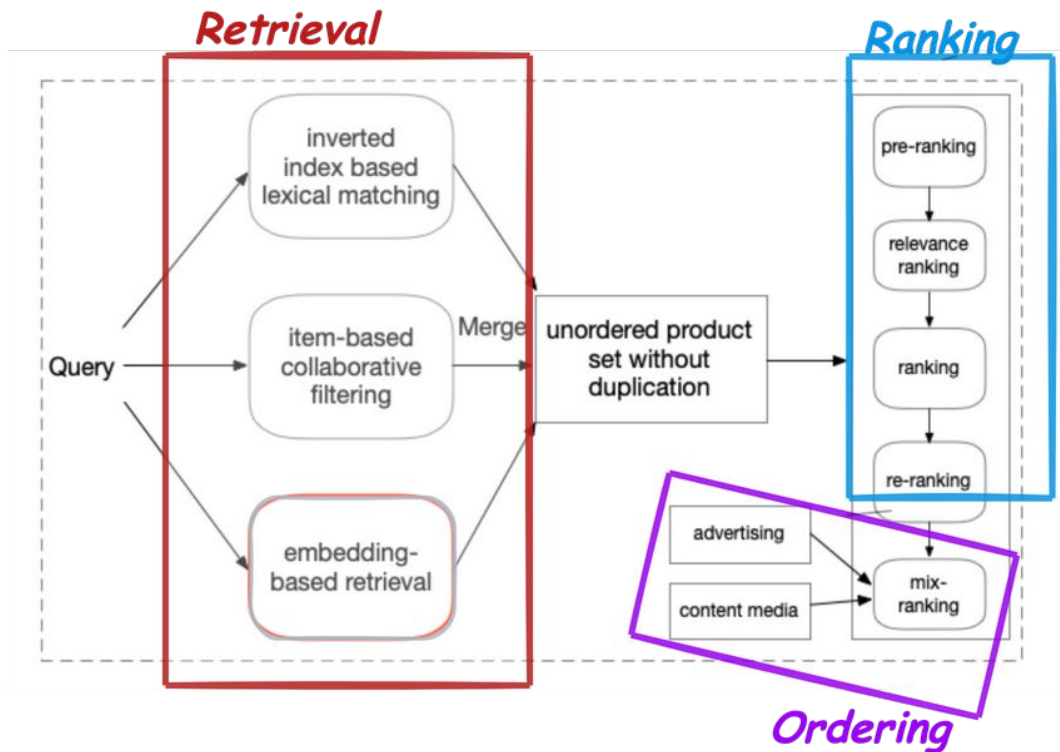
In 2019, Instagram Explore feed was made of

- Sparse Retriever: 500 candidates
- Lightweight NN: 150 -> 50
- Deep Neural Network: 50 -> 25

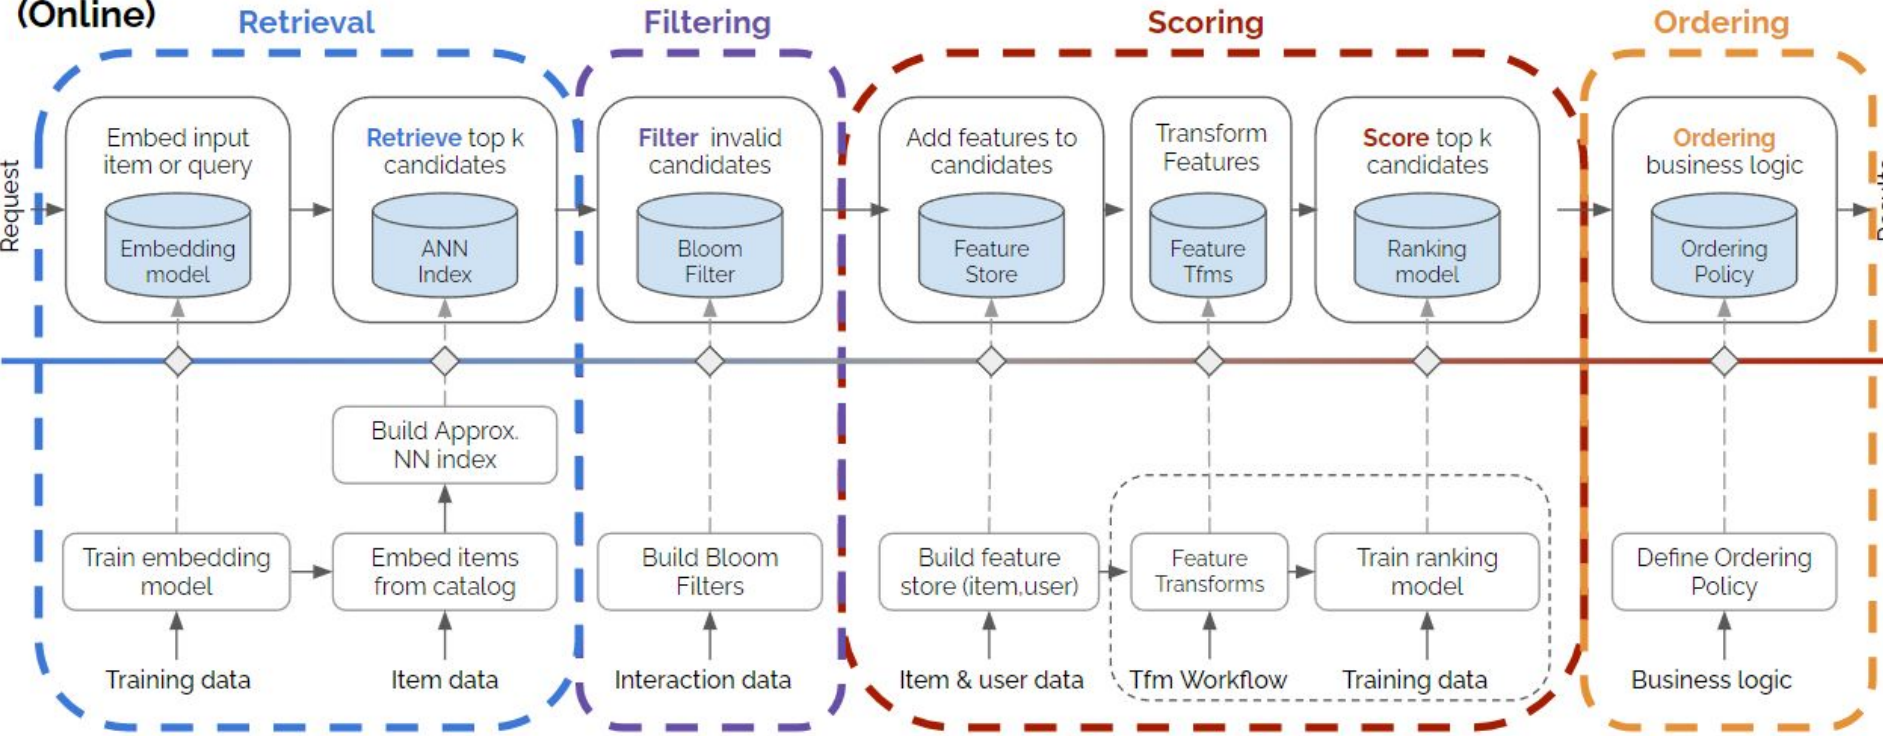
Medvedev, Ivan, Haotian Wu, and Taylor Gordon. "Powered by AI: Instagram's Explore Recommender System." Powered by AI: Instagram's Explore recommender system, November 2019.

<https://ai.facebook.com/blog/powered-by-ai-instagrams-explore-recommender-system/>

Overview of Taobao Search




**Inference
(Online)**



**Training
(Offline)**

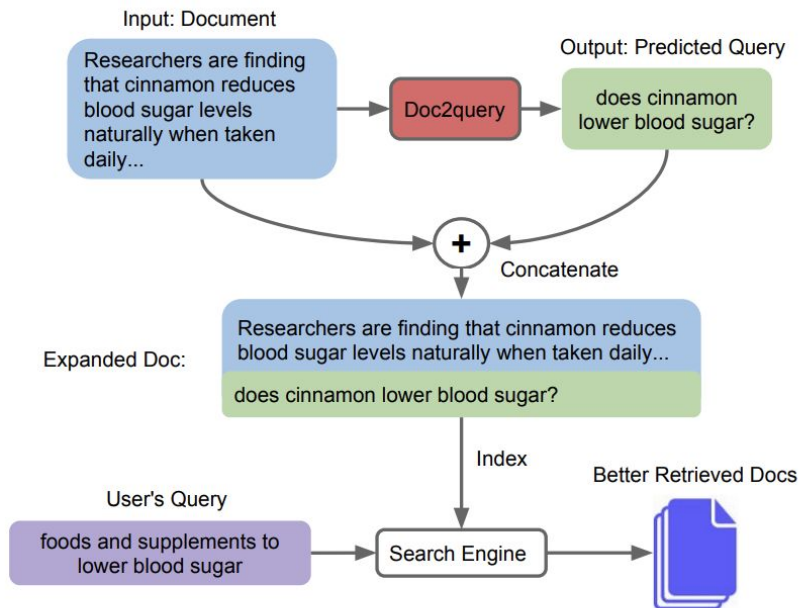
| | Retrieval | Filtering | Scoring | Ordering |
|-------------------|--|---|--|---|
| Music Discovery | Find similar songs based on nearest neighbour search | Remove tracks users listened before | Predict likelihood that a user will listen to a song | Trade-Off between score, similarity, BPM, etc |
| Social Media | Find new posts in user's network | Remove posts from blocked and muted users | Predict likelihood that a user will interact with it | Change order that adjust posts are from different authors |
| Online Store | Find items which are usually co-purchased | Remove items which are out of stock | Predict likelihood that a user will purchase an item | Reorder items based on price points |
| Streaming Service | Find items based on different rows/shelves/topics | Remove items which are not available for user's country | Predict user's stream time per item | Organize recommendations to fit genre distributions |

Zipfian Distribution of Web Traffic



- Most web platforms traffic follows a zipfian distribution
- Frequent queries / head traffic can be cached
-

Enhancing Sparse Index: Doc2Query / Doc2T5Query



- Use a causal language model to generate additional text to add to documents when indexing.
- At retrieval time, use BM25

Benchmarking

— — —

| Dataset | BM25 | Dense (TAS-B) | | BM25 on docT5query | BM25 + Cross Encoder |
|-------------------------|-------|------------------|---|-----------------------|----------------------------|
| Baseline (MS MARCO) | 0.228 | 0.408 | — | <u>0.338</u> | <u>0.413</u> |
| Quora | 0.789 | 0.835 | — | <u>0.802</u> | <u>0.825</u> |
| DBPedia | 0.313 | 0.384 | — | <u>0.331</u> | <u>0.409</u> |
| — | — | — | — | — | — |
| TREC-COVID (Medical) | 0.656 | 0.481 | — | <u>0.713</u> | <u>0.757</u> |
| Signal-1M (Tweets) | 0.330 | 0.289 | — | <u>0.325</u> | <u>0.338</u> |

Do models trained on MS MARCO work for different datasets ?

Dense Retrieval performs well on similar domain

COVID/Tweets are out of domain

BM25 + DR perform best

Cost to Serve

— — —

| Model | Dimension | Latency (CPU) | Latency (GPU) | Index Size |
|--------------------------------------|-----------|---------------|---------------|------------|
| BM-25 | – | 20 ms | – | 0.4 GB |
| docT5Query | – | 30 ms | – | 0.4 GB |
| Dense Embedding Passage Retrieval | 768 | 230 ms | 19 ms | 3 GB |
| BM25 + Cross Encoder | – | 6100 ms | 450 ms | 0.4 GB |

Estimated average
retrieval latency and
index sizes

Dataset: DBPedia (1
million docs)

Lower Latency and
Index Size is
preferred

CPU: 8 core Intel
Xeon Platinum 8168
CPU @ 2.70GHz

GPU: 1 Nvidia Tesla
V100

Resources

— — —

- [Faiss Missing Manual \(Pinecone\)](#)
- [Natural Language Processing \(NLP\) for Semantic Search \(Pinecone\)](#)
- [CIKM 2021 Tutorial: IR From Bag-of-words to BERT](#)
- [Haystack US 2021 - Semantic Product Search - Vector Search for E-Commerce - Simon Hughes](#)